

Data Wrangling in R

Erica Bishop

2023-02-24

Data Wrangling in R Examples

The code below is excerpted from a machine learning lab in which my goal was to build four different types of classification models to predict whether a song belonged to my saved track list on Spotify or my lab partner's saved track list. The original work is saved in the `labg5_get_spotify_data.R` script and the `lab5.Rmd` markdown file within this repository if you'd like to check out the resulting models and accuracy comparisons. Metadata and variable definitions are available [here](#).

The spotify data wrangling presented a unique challenge because the API requests are limited to 50 rows at a time, so although the data was pretty clean out of the gate in terms of missing data and variable consistency, it required some manipulation to join and process for modeling. I've included a brief demonstration of how I've used the `janitor` and `lubridate` packages to work with data that was messy in other ways at the end of this document

Preparing API data for classification models

Step 1: Request my data from the Spotify API

```
###Get access token
#access and edit .renviron with usethis::edit_r_environ
access_token <- get_spotify_access_token(client_id = Sys.getenv("SPOTIFY_CLIENT_ID"), #client secret and
                                         client_secret = Sys.getenv("SPOTIFY_CLIENT_SECRET"))

###Request saved track list
offsets_list <- c(seq(0, 2900, 50)) #create list to get saved songs

saved_tracks_eb <- lapply(X = offsets_list,
                          FUN = get_my_saved_tracks,
                          limit = 50 #produce list of dfs for each call of 50
) |>
  bind_rows() #bind lists into one df

#Request saved track attributes
track_id_list <- saved_tracks_eb$track.id #create list of track ids

track_attributes_eb <- lapply(X = track_id_list, #get audio features for each track id
                              FUN = get_track_audio_features) |>
  bind_rows() #combine resulting list of dfs
```

Step 2: Wrangle my data into a usable format for my lab partner (Jillian)

```

#append track list name to the track attributes df
saved_track_names <- saved_tracks_eb |>
  select(track.name, track.id) |> #select out just track name and id
  rename(id = track.id) #rename id column to join

eb_tracks_df <- left_join(track_attributes_eb, saved_track_names, by = "id") #join df

#add user column
eb_tracks_df <- eb_tracks_df |>
  add_column(user = "Erica")

#save dataframe as a csv to send to Jillian
write_csv(eb_tracks_df, here("eb_track_attributes.csv"))

```

Step 3: Load my lab partner's data

```

#read in Jillian's data
ja_tracks_df <- read_csv(here("jillian_spotify_data.csv")) |> #saved to repo for ease of re-running
  select(-track_id) |> #drop track_id column (repeats id column)
  rename(track.name = track_name) #renametrack_name to match

## Rows: 931 Columns: 21
## -- Column specification -----
## Delimiter: ","
## chr (8): type, id, uri, track_href, analysis_url, track_name, track_id, user
## dbl (13): danceability, energy, key, loudness, mode, speechiness, acousticne...
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.

```

Step 4: Join data and investigate overlaps

For this lab, I decided to create a third classification for songs that were in both my and my partner's saved track lists.

```

#take a closer look at variables and see if the names match
names(eb_tracks_df)

```

```

## [1] "danceability" "energy" "key" "loudness"
## [5] "mode" "speechiness" "acousticness" "instrumentalness"
## [9] "liveness" "valence" "tempo" "type"
## [13] "id" "uri" "track_href" "analysis_url"
## [17] "duration_ms" "time_signature" "track.name" "user"

```

```
names(ja_tracks_df)
```

```

## [1] "danceability" "energy" "key" "loudness"
## [5] "mode" "speechiness" "acousticness" "instrumentalness"
## [9] "liveness" "valence" "tempo" "type"
## [13] "id" "uri" "track_href" "analysis_url"
## [17] "duration_ms" "time_signature" "track.name" "user"

```

```
#Investigate the overlap of saved songs
```

```

eb_ja_shared_faves <- semi_join(eb_tracks_df, ja_tracks_df, by = "id") #use semi_join to get df of tracks
#Over 100 tracks in common - worth creating a third class

```

```

#Combine both datasets
eb_ja_tracks <- full_join(eb_tracks_df,
                          ja_tracks_df,
                          by = c("danceability", "energy", "key", "loudness", #joining by all variables
                                "mode", "speechiness", "acousticness", "instrumentalness",
                                "liveness", "valence", "tempo", "type", "id", "uri", "track_href",
                                "analysis_url", "duration_ms", "time_signature", "track.name")) |>

mutate(user = case_when(
  user.y == "Jillian" & user.x == "Erica" ~ "Both", #relabel songs in both user lists with case_when
  user.y == "Jillian" & is.na(user.x) ~ "Jillian",
  TRUE ~ "Erica"
)) |> #drop individual user columns and a few other less useful predictors
select(-c(user.x, user.y, uri, track_href, analysis_url, type))

#save output df as a csv
write_csv(eb_ja_tracks, here("eb_ja_tracks.csv"))

```

Step 5: Pre-process data for modeling (encoding and standardizing)

```

#remove unique identifiers from dataset (track id and song name)
tracks_df <- eb_ja_tracks |>
  select(-c(id, track.name)) |>
  mutate(user = as.factor(user)) #make outcome variable a factor

#split the data FIRST
set.seed(123) #set seed for reproducibility

tracks_split <- initial_split(tracks_df) #split
tracks_train <- training(tracks_split) #training dataset
tracks_test <- testing(tracks_split) #testing dataset

#Preprocess the data for tidymodels modeling - encoding with recipe
tracks_recipe <- recipe(user ~., data = tracks_train) |>
  step_dummy(all_nominal(), -all_outcomes(), one_hot = TRUE) |> #one hot encoding of nominal variables
  step_normalize(all_numeric(), -all_outcomes()) |> #normalize scale of numeric variables
  prep()

```

Cleaning messy data with Janitor, Lubridate, and zoo

For my final project in my statistics course, I wanted to examine the relationship between the number of electric vehicles registered and the ground-level ozone pollution in the Denver Metro area. A full write-up on this project can be accessed through my blog post.

The EV data was downloaded from <https://www.atlasevhub.com/materials/state-ev-registration-data/#data> - zip codes are for some reason in multiple different formats and lengths - over 1300 unique zips (should be no more than 500 or so) - date range: January 2010 - July 2022 - includes fuel cell, plug in hybrid, and EV - NOT using cumulative counts because this dataset includes original registrations and renewals without distinction, so according to the recommendations from the data source, treating like a snapshot in time

The air quality data is from the EPA's Outdoor Air Quality Data, downloadable here. - Includes Daily AQI, Daily Ozone, and Annual AQI data by county

```

#set up data directory
datadir <- ("/Users/ericabishop/Documents/MEDS-fall-classwork/EDS222-stats/final_project/data") #contai

```

```
figdir <- ("/Users/ericabishop/Documents/MEDS-fall-classwork/EDS222-stats/final_project/figs")#to save
```

Step 1: read-in and wrangle EV registration and county/zip code data

EV Data

#read in EV data and format dates

```
co_evs <- read_csv(file.path(datadir, "co_ev_registrations_public.csv")) |>
  janitor::clean_names()
```

```
## Rows: 1246097 Columns: 10
```

```
## -- Column specification -----
```

```
## Delimiter: ","
```

```
## chr (8): DMV Snapshot, State, Registration Valid Date, Registration Expirati...
```

```
## dbl (2): DMV ID, ZIP Code
```

```
##
```

```
## i Use `spec()` to retrieve the full column specification for this data.
```

```
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

#format dates

```
co_evs <- co_evs |>
  mutate(registration_valid_date = lubridate::mdy(registration_valid_date),
         registration_expiration_date = lubridate::mdy(registration_expiration_date))
```

#read in county data to validate zips and match with EPA data, drop state column

```
co_counties <- read_csv(file.path(datadir, "CO_counties.csv")) |>
  clean_names() |>
  select("county", "zip_code") #validating zip codes by those that match a county
```

```
## Rows: 661 Columns: 3
```

```
## -- Column specification -----
```

```
## Delimiter: ","
```

```
## chr (2): State, County
```

```
## dbl (1): ZIP Code
```

```
##
```

```
## i Use `spec()` to retrieve the full column specification for this data.
```

```
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

#attach county names associated with zip code to co_evs

```
co_evs_counties <- left_join(co_evs, co_counties, by = "zip_code")
```

#rename date column to join

```
co_evs_counties <- co_evs_counties |>
  rename(date_local = registration_valid_date)
```

#manipulate variables of interest into one DF - BY COUNTY

#select columns of interest from EVs into new DF

```
evs_clean <- co_evs_counties |>
  select("county", "date_local", "zip_code") |>
  mutate(my_date = as.yearmon(date_local))
```

Step 2: Read in EPA ozone data and filter to Colorado

OZONE Data

#read in daily ozone data

```
daily_ozone_2010 <- read_csv(file.path(datadir, "daily_summary_ozone", "daily_44201_2010.csv"))
```

```
## Rows: 378545 Columns: 29
```

```
## -- Column specification -----
```

```
## Delimiter: ","
```

```
## chr (16): State Code, County Code, Site Num, Datum, Parameter Name, Sample ...
```

```
## dbl (10): Parameter Code, POC, Latitude, Longitude, Observation Count, Obse...
```

```
## lgl (1): Method Code
```

```
## date (2): Date Local, Date of Last Change
```

```
##
```

```
## i Use `spec()` to retrieve the full column specification for this data.
```

```
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
daily_ozone_2011 <- read_csv(file.path(datadir, "daily_summary_ozone", "daily_44201_2011.csv"))
```

```
## Rows: 383872 Columns: 29
```

```
## -- Column specification -----
```

```
## Delimiter: ","
```

```
## chr (16): State Code, County Code, Site Num, Datum, Parameter Name, Sample ...
```

```
## dbl (10): Parameter Code, POC, Latitude, Longitude, Observation Count, Obse...
```

```
## lgl (1): Method Code
```

```
## date (2): Date Local, Date of Last Change
```

```
##
```

```
## i Use `spec()` to retrieve the full column specification for this data.
```

```
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
daily_ozone_2012 <- read_csv(file.path(datadir, "daily_summary_ozone", "daily_44201_2012.csv"))
```

```
## Rows: 386439 Columns: 29
```

```
## -- Column specification -----
```

```
## Delimiter: ","
```

```
## chr (16): State Code, County Code, Site Num, Datum, Parameter Name, Sample ...
```

```
## dbl (10): Parameter Code, POC, Latitude, Longitude, Observation Count, Obse...
```

```
## lgl (1): Method Code
```

```
## date (2): Date Local, Date of Last Change
```

```
##
```

```
## i Use `spec()` to retrieve the full column specification for this data.
```

```
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
daily_ozone_2013 <- read_csv(file.path(datadir, "daily_summary_ozone", "daily_44201_2013.csv"))
```

```
## Rows: 389373 Columns: 29
```

```
## -- Column specification -----
```

```
## Delimiter: ","
```

```
## chr (16): State Code, County Code, Site Num, Datum, Parameter Name, Sample ...
```

```
## dbl (10): Parameter Code, POC, Latitude, Longitude, Observation Count, Obse...
```

```
## lgl (1): Method Code
```

```
## date (2): Date Local, Date of Last Change
```

```
##
```

```
## i Use `spec()` to retrieve the full column specification for this data.
```

```
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```

daily_ozone_2014 <- read_csv(file.path(datadir, "daily_summary_ozone", "daily_44201_2014.csv"))

## Rows: 388613 Columns: 29
## -- Column specification -----
## Delimiter: ","
## chr  (16): State Code, County Code, Site Num, Datum, Parameter Name, Sample ...
## dbl  (10): Parameter Code, POC, Latitude, Longitude, Observation Count, Obse...
## lgl   (1): Method Code
## date  (2): Date Local, Date of Last Change
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
daily_ozone_2015 <- read_csv(file.path(datadir, "daily_summary_ozone", "daily_44201_2015.csv"))

## Rows: 386404 Columns: 29
## -- Column specification -----
## Delimiter: ","
## chr  (16): State Code, County Code, Site Num, Datum, Parameter Name, Sample ...
## dbl  (10): Parameter Code, POC, Latitude, Longitude, Observation Count, Obse...
## lgl   (1): Method Code
## date  (2): Date Local, Date of Last Change
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
daily_ozone_2016 <- read_csv(file.path(datadir, "daily_summary_ozone", "daily_44201_2016.csv"))

## Rows: 392952 Columns: 29
## -- Column specification -----
## Delimiter: ","
## chr  (16): State Code, County Code, Site Num, Datum, Parameter Name, Sample ...
## dbl  (10): Parameter Code, POC, Latitude, Longitude, Observation Count, Obse...
## lgl   (1): Method Code
## date  (2): Date Local, Date of Last Change
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
daily_ozone_2017 <- read_csv(file.path(datadir, "daily_summary_ozone", "daily_44201_2017.csv"))

## Rows: 404217 Columns: 29
## -- Column specification -----
## Delimiter: ","
## chr  (16): State Code, County Code, Site Num, Datum, Parameter Name, Sample ...
## dbl  (10): Parameter Code, POC, Latitude, Longitude, Observation Count, Obse...
## lgl   (1): Method Code
## date  (2): Date Local, Date of Last Change
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
daily_ozone_2018 <- read_csv(file.path(datadir, "daily_summary_ozone", "daily_44201_2018.csv"))

## Rows: 402906 Columns: 29
## -- Column specification -----
## Delimiter: ","

```

```

## chr (16): State Code, County Code, Site Num, Datum, Parameter Name, Sample ...
## dbl (10): Parameter Code, POC, Latitude, Longitude, Observation Count, Obse...
## lgl (1): Method Code
## date (2): Date Local, Date of Last Change
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
daily_ozone_2019 <- read_csv(file.path(datadir, "daily_summary_ozone", "daily_44201_2019.csv"))

## Rows: 389651 Columns: 29
## -- Column specification -----
## Delimiter: ","
## chr (16): State Code, County Code, Site Num, Datum, Parameter Name, Sample ...
## dbl (10): Parameter Code, POC, Latitude, Longitude, Observation Count, Obse...
## lgl (1): Method Code
## date (2): Date Local, Date of Last Change
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
daily_ozone_2020 <- read_csv(file.path(datadir, "daily_summary_ozone", "daily_44201_2020.csv"))

## Rows: 391845 Columns: 29
## -- Column specification -----
## Delimiter: ","
## chr (16): State Code, County Code, Site Num, Datum, Parameter Name, Sample ...
## dbl (10): Parameter Code, POC, Latitude, Longitude, Observation Count, Obse...
## lgl (1): Method Code
## date (2): Date Local, Date of Last Change
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
daily_ozone_2021 <- read_csv(file.path(datadir, "daily_summary_ozone", "daily_44201_2021.csv"))

## Rows: 390562 Columns: 29
## -- Column specification -----
## Delimiter: ","
## chr (16): State Code, County Code, Site Num, Datum, Parameter Name, Sample ...
## dbl (10): Parameter Code, POC, Latitude, Longitude, Observation Count, Obse...
## lgl (1): Method Code
## date (2): Date Local, Date of Last Change
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
daily_ozone_2022 <- read_csv(file.path(datadir, "daily_summary_ozone", "daily_44201_2022.csv"))

## Rows: 233601 Columns: 29
## -- Column specification -----
## Delimiter: ","
## chr (16): State Code, County Code, Site Num, Datum, Parameter Name, Sample ...
## dbl (10): Parameter Code, POC, Latitude, Longitude, Observation Count, Obse...
## lgl (1): Method Code
## date (2): Date Local, Date of Last Change
##

```

```
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
#write function to filter data to just Colorado
```

```
filter_CO <- function(df) {
  df_filtered <- df |>
    filter(`State Name` == "Colorado") |>
    clean_names()
  return(df_filtered)
}
```

```
#filter to just CO
```

```
dozo_10_CO <- filter_CO(daily_ozone_2010)
dozo_11_CO <- filter_CO(daily_ozone_2011)
dozo_12_CO <- filter_CO(daily_ozone_2012)
dozo_13_CO <- filter_CO(daily_ozone_2013)
dozo_14_CO <- filter_CO(daily_ozone_2014)
dozo_15_CO <- filter_CO(daily_ozone_2015)
dozo_16_CO <- filter_CO(daily_ozone_2016)
dozo_17_CO <- filter_CO(daily_ozone_2017)
dozo_18_CO <- filter_CO(daily_ozone_2018)
dozo_19_CO <- filter_CO(daily_ozone_2019)
dozo_20_CO <- filter_CO(daily_ozone_2020)
dozo_21_CO <- filter_CO(daily_ozone_2021)
dozo_22_CO <- filter_CO(daily_ozone_2022)
```

```
#combine into one df with rbind
```

```
daily_ozo_CO <- rbind(dozo_10_CO, dozo_11_CO, dozo_12_CO, dozo_13_CO, dozo_14_CO, dozo_15_CO, dozo_16_CO,
```

```
#remove precursor datasets from memory to free up space
```

```
rm(dozo_10_CO, dozo_11_CO, dozo_12_CO, dozo_13_CO, dozo_14_CO, dozo_15_CO, dozo_16_CO, dozo_17_CO, dozo_18_CO, dozo_19_CO, dozo_20_CO, dozo_21_CO, dozo_22_CO)
rm(daily_ozone_2010, daily_ozone_2011, daily_ozone_2012, daily_ozone_2013, daily_ozone_2014, daily_ozone_2015, daily_ozone_2016, daily_ozone_2017, daily_ozone_2018, daily_ozone_2019, daily_ozone_2020, daily_ozone_2021, daily_ozone_2022)
```

Step 3: Combine data and aggregate to monthly time scale

```
#select columns of interest from ozone into new DF
```

```
ozone_clean <- daily_ozo_CO |>
  select("date_local", "units_of_measure", "parameter_name", "aqi", "county_name", "local_site_name", "county_fips") |>
  mutate(my_date = as.yearmon(date_local)) |>
  rename(county = county_name)
```

```
#aggregate both dfs monthly for analysis
```

```
evs_summary <- evs_clean |>
  group_by(my_date) |>
  summarize(ev_reg_count = n())
```

```
ozone_summary <- ozone_clean |>
  group_by(my_date) |>
  summarize(monthly_ozone_avg_ppm = mean(arithmetic_mean))
```

```
#combine summary tables into one dataframe of monthly/yearly (my) data
```

```
my_summary <- left_join(evs_summary, ozone_summary, by = "my_date")
```